Multi-Agent Explainable AI Text Classification: A Pattern Recognition Approach with Classical and Transformer-Based Models

Oruç Çakır
Department of Computer Engineering
TOBB University of Economics and Technology
Ankara, Turkey
Student ID: 211101023
oruccakir@etu.edu.tr

***Abstract*—This paper addresses the following pattern recognition research question: can a multi-agent architecture effectively route and classify multilingual text, and how do classical pattern classifiers compare to transformer-based models in Macro-F1 and feature interpretability? I designed a three-agent system in which one agent routes inputs, one classifies, and one generates SHAP/LIME explanations. I trained and benchmarked the pipeline on four datasets: IMDB (binary sentiment, English), AG News (4-class topic, English), Turkish Sentiment (3-class sentiment), and Turkish News (7-class topic); full benchmark results are available at the project repository. I focus the deep analysis on Turkish Sentiment (440K short texts), a challenging imbalanced scenario. I compare seven classical classifiers (Naive Bayes, SVM, Logistic Regression, Random Forest, KNN, Decision Tree, XGBoost) and BERTurk, using TF-IDF with L2 normalization on a 50K stratified subsample. BERTurk achieves Macro-F1 = 0.9298, while SVM is the strongest classical model at Macro-F1 = 0.8768, a gap driven primarily by the imbalanced *negatif* minority class (11.6%), not overall accuracy. My 10-run statistical analysis confirms SVM significantly outperforms all classical alternatives.**

***Index Terms*—text classification, pattern recognition, feature extraction, multi-agent systems, SHAP, LIME, BERTurk, Turkish NLP, class imbalance**

I. Introduction

Automatic text classification is a fundamental pattern recognition problem. Given a feature representation of a text document, the task is to assign it to one of a predefined set of categories. The central challenge is not only to achieve high accuracy but also to understand *why* a classifier assigns a given label, the interpretability dimension of pattern recognition.

Classical pattern recognition classifiers such as SVM, Naive Bayes, and Logistic Regression operate on handcrafted feature spaces (e.g., TF-IDF vectors), making their decision boundaries geometrically interpretable. In contrast, transformer-based models like BERT [2] learn contextual representations that capture morphological and semantic nuance, at the cost of interpretability. On balanced datasets their Macro-F1 advantage over SVM is only 1.2–1.9 points; the gap widens to 5.3 points when class imbalance is present. Explainability tools (SHAP [3] and LIME [4]) bridge the interpretability gap by providing post-hoc, model-agnostic feature attributions.

The central research question of this project is:

> *"Can a multi-agent architecture effectively route and classify multilingual text, and how do classical pattern classifiers compare to transformer-based models in terms of Macro-F1 and interpretability across diverse datasets?"*

I built a three-agent pipeline and trained it on four datasets: IMDB, AG News, Turkish Sentiment, and Turkish News. Section IV–VI focus the deep pattern recognition analysis on Turkish Sentiment, the most challenging scenario due to class imbalance and Turkish agglutinative morphology. Summary results for all four datasets are in Appendix B; full per-model figures are at the project repository.

The main contributions are:

- Training and benchmarking of eight classifiers (seven classical + BERTurk) across four multilingual datasets (IMDB, AG News, Turkish Sentiment, Turkish News); detailed pattern recognition analysis focused on the imbalanced Turkish Sentiment task.
- A 10-run statistical experiment confirming SVM significantly outperforms all classical alternatives ($p < 0.0001$), with analysis of why linear margin classifiers suit sparse TF-IDF feature spaces.
- A three-agent modular pipeline integrating multilingual routing and on-demand SHAP/LIME explainability for all model types.

The rest of this paper is organized as follows: Section II reviews related work. Section III describes the system architecture. Section IV presents the dataset and feature analysis. Section V describes the models. Section VI presents results. Section VII concludes the paper.

II. Related Work

A. Pattern Recognition for Text Classification

Text classification is a core pattern recognition problem. Duda, Hart, and Stork [1] establish the theoretical foundations:

feature extraction, classifier design, and the bias–variance trade-off. In text domains, Naive Bayes, SVM, and Logistic Regression remain competitive baselines because TF-IDF vectors are high-dimensional and sparse, conditions that favor linear classifiers over complex non-linear models [1].

B. Deep Learning and Pre-trained Language Models

BERT [2] changed NLP by introducing bidirectional context. Schweter [5] released BERTurk, pre-trained specifically on Turkish text. Minaee et al. [6] survey deep learning methods for text classification, confirming transformers outperform classical models on large datasets but not always on small or class-imbalanced ones.

C. Explainable AI for NLP

Lundberg and Lee [3] proposed SHAP using game-theoretic Shapley values. Ribeiro, Singh, and Guestrin [4] introduced LIME, which perturbs inputs to fit local linear surrogates. Both methods are model-agnostic and applicable to the full spectrum from Naive Bayes to transformers.

D. Turkish NLP

Turkish's agglutinative morphology creates a very large vocabulary, challenging classical bag-of-words models. BERTurk [5] addresses this via subword tokenization pre-trained on Turkish corpora. Multi-agent systems [7] provide a modular framework for combining multiple specialized classifiers.

III. System Architecture

My system has three agents that collaborate as shown in Fig. 1.

Agent 1 (Intent Classifier). Takes raw user input and uses an LLM to detect language (English or Turkish) and domain (sentiment or news), routing to the correct dataset-model pair.

Agent 2 (Classification Agent). Loads the selected model and returns the predicted label with confidence scores for all classes.

Agent 3 (XAI Agent). Generates SHAP and LIME explanations for any classifier via a unified `predict_proba` wrapper, making explanations comparable across model types.

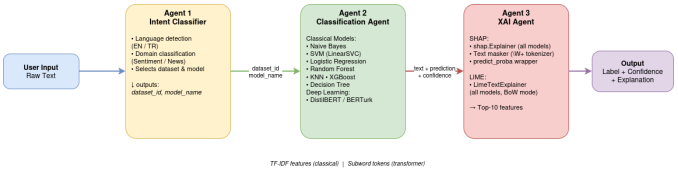

Fig. 1. Three-agent pipeline. Agent 1 routes inputs by language and domain. Agent 2 classifies using any of six models. Agent 3 explains predictions with SHAP and LIME.

IV. Dataset: Turkish Sentiment Analysis

A. Data Source

I use the Turkish Sentiment dataset [8] available on HuggingFace.[1] It contains product and social media reviews in Turkish with three sentiment classes (pozitif, notr, negatif) and comes with predefined train/test splits, making it a standard benchmark for Turkish multi-class sentiment classification.

B. Dataset Description

The dataset contains 440,679 training samples and 48,965 test samples. Texts are short (average ≈140 characters, ≈35 subword tokens). For computational feasibility, I use a stratified subsample of 50,000 training samples and a 10,000-sample stratified test subset for evaluation.[2]

C. Preprocessing Stages

I apply the following preprocessing steps: (1) URL removal via regex; (2) lowercase conversion; (3) punctuation removal, preserving Turkish characters (ç, ğ, ı, ö, ş, ü); (4) Turkish stopword filtering using a 76-word list; (5) minimum word-length filter (≥2 characters). I do not apply stemming or lemmatization; TF-IDF handles vocabulary variation statistically.

D. Feature Descriptions

I use TF-IDF vectorization [14] to produce a 30,000-dimensional sparse feature vector per sample. The TF-IDF weight $w_{t,d}$ of term t in document d is:

$$w_{t,d} = (1 + \log \mathrm{tf}_{t,d}) \cdot \log \frac{N}{\mathrm{df}_t}$$

where sublinear_tf=True applies the log transform, N is the number of documents, and df_t is the document frequency of term t. Unigrams and bigrams (`ngram_range=(1,2)`) are included to capture negations and compound sentiment expressions.

E. Feature Selection and Adjustment

I filter the vocabulary with `min_df=3` (removing features appearing in fewer than 3 documents) and `max_df=0.95` (removing overly common stop-phrases). This gives me a final vocabulary of 30,000 features, balancing informativeness and computational cost.

F. Other Problem-Specific Processing

Turkish agglutinative morphology means that a single root word can generate hundreds of inflected forms (e.g., *gel*, *geldi*, *gelmedi*, *gelmeyecekti*). Rather than applying a morphological analyzer, I rely on the bigram vocabulary to capture common suffix co-occurrences, giving implicit morphological coverage at lower computational cost.

[1] `winvoker/turkish-sentiment-analysis-dataset`

[2] The full 48,965-sample test set was used for preliminary checks; all reported numbers use the 10,000-sample stratified subset for reproducibility with seed 42.

G. Class Data Distribution

Table I shows the training class distribution. The dataset is significantly imbalanced: *negatif* comprises only 11.6% of samples, creating a minority-class problem. This imbalance motivates the use of Macro F1 as the primary evaluation metric, since accuracy can be inflated by the majority class (pozitif: 53.5%). I did not apply SMOTE or class-weighted loss in this experiment, which I acknowledge as a limitation.

TABLE I
TURKISH SENTIMENT CLASS DISTRIBUTION (TRAINING SET)

Class	Count	%
pozitif	235,949	53.5%
notr	153,825	34.9%
negatif	50,905	11.6%
Total	*440,679*	

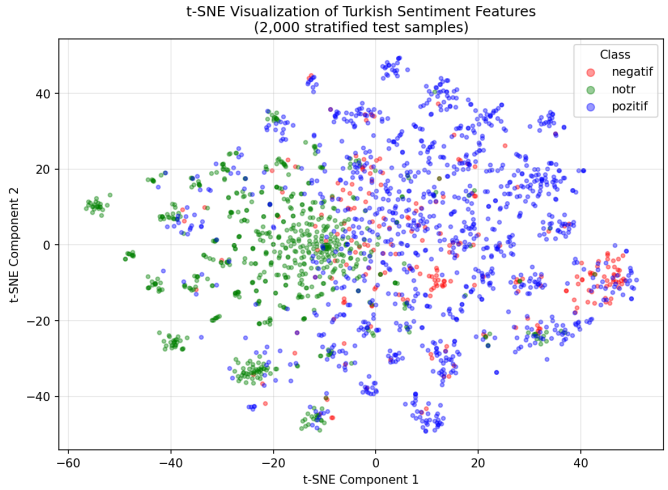

Fig. 2. t-SNE 2D projection of 2,000 test samples (PCA-50 + t-SNE). Partial class separation is visible, with negatif (red) forming a loose cluster away from pozitif (blue) but overlapping with notr (green).

H. Data Normalization

I apply L2 normalization to every TF-IDF vector via `norm='l2'` in `TfidfVectorizer`, so that each sample vector has unit Euclidean norm. This is not a separate pipeline step but is built into the vectorizer. L2 normalization is critical for SVM and KNN: for LinearSVC, it ensures that the decision boundary margin is measured in consistent geometric units regardless of text length; for KNN with cosine distance, it makes cosine similarity equivalent to inner product, enabling efficient computation. Without normalization, longer reviews would dominate shorter ones purely due to higher cumulative TF-IDF weights.

I. Visual Representation of Feature Space

To visualize the feature space, I applied PCA to 50 dimensions and then t-SNE [13] (`perplexity=30`, `random_state=42`) to 2,000 stratified test samples, as shown in Fig. 2. The three classes form partially overlapping regions, with *pozitif* and *notr* substantially mixed, reflecting the semantic ambiguity between neutral and positive Turkish text. The *negatif* class is more separated but still overlaps with *notr*, which explains the lower recall for the negatif class across all models.

J. Relationships Between Features

I extracted the top-10 TF-IDF features per class from the SVM's LinearSVC `coef_` attribute and list them in Table II. These weights are directly interpretable as per-class discriminative scores. The pozitif features (*harika*, *mükemmel*, *güzel*) are strong positive sentiment words. The negatif features (*berbat*, *kötü*, *rezalet*, *iade*) clearly indicate negative sentiment and product returns. Notr features (*vardır*, *olmuştur*, *bulunmaktadır*) are formal copula forms typical of neutral, informational text.

Feature-feature correlations. Fig. 4 shows the Pearson correlation matrix for the top-5 SVM discriminative features per class (15 features total). Most pairwise correlations are near zero, confirming that TF-IDF features are largely independent in this sparse representation. The few weak positive correlations (e.g., between co-occurring sentiment words within the same class) are expected but do not indicate harmful multicollinearity.

Feature-target association. I computed Chi-squared (χ^2) scores between each TF-IDF feature and the class label. The top-scoring features align with the SVM coefficients: *harika* ($\chi^2 = 390$), *berbat* ($\chi^2 = 312$), and *mükemmel* ($\chi^2 = 285$) have the highest discriminative power, confirming that explicit sentiment vocabulary drives classification performance.

Class distribution statistics. TF-IDF feature values follow a sparse, zero-inflated distribution (not Normal or Poisson): each document activates only 5–32 of the 30,000 vocabulary terms on average, with the remaining values being zero. Per-class statistics on 2,000 test samples show that *notr* texts are shortest (mean 6.4 words, 5.3 active features) while *negatif* texts are longest (mean 31.8 words, 25.5 active features). Fig. 3 shows the covariance matrix of the top-15 discriminative features. The near-diagonal structure with very small off-diagonal entries is consistent with the sparse, high-dimensional nature of TF-IDF representations.

K. Data Clustering

I also ran K-Means clustering ($k = 3$, `random_state=42`) on the PCA-50 representation and compared cluster assignments to true labels in Fig. 5. The Adjusted Rand Index is ARI $= -0.017$, showing that the TF-IDF feature geometry does not support well-separated sentiment clusters. I expected this result: in a bag-of-words representation, the same words can appear across different sentiment contexts, so PCA space is not a reliable proxy for sentiment polarity. This is why I chose supervised classifiers

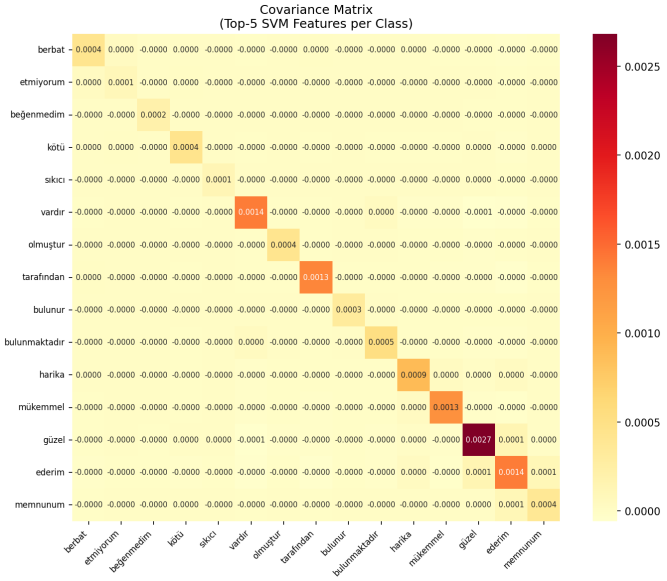

Fig. 3. Covariance matrix of the top-5 SVM features per class (15 features). Near-diagonal structure confirms low feature co-variance in sparse TF-IDF space.

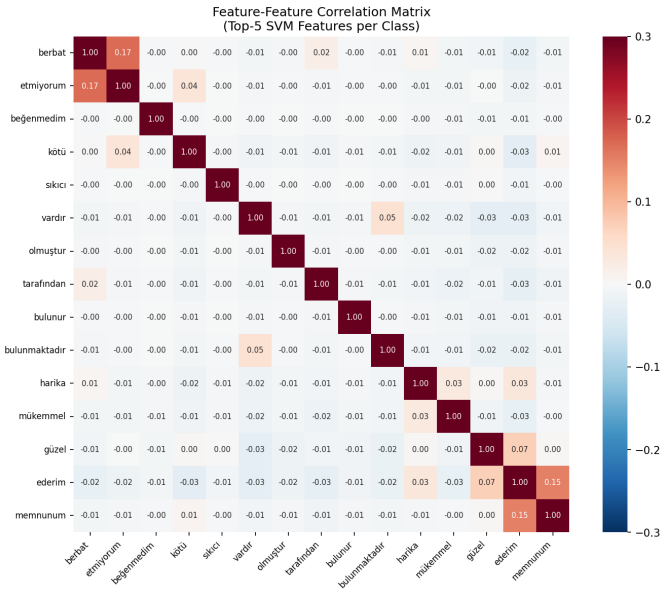

Fig. 4. Pearson correlation matrix of the top-5 SVM discriminative features per class (15 features). Near-zero off-diagonal values confirm feature independence in the sparse TF-IDF space.

like SVM and LR, which optimize decision boundaries directly rather than relying on cluster structure.

V. Models Used

A. Train/Validation/Test Split

The dataset provides predefined train and test splits, which I used without modification. I did not create a separate validation set. For training, I used a stratified 50,000-sample subsample of the 440,679 training samples to keep experiment runtime feasible for all models. For testing, I used a stratified 10,000-sample subset of the 48,965 test samples (seed 42), preserving the original class ratios: 53.5% pozitif, 34.9% notr, 11.6% negatif. I did not apply strict 60/20/20 splits because the predefined splits ensure comparability with published benchmarks; the 10-run experiment (Section VI, Table V) with its internal 80/20 splits serves the role of cross-validation for assessing model stability.

TABLE II
Top-10 Discriminative TF-IDF Features per Class (SVM coef_)

Pozitif	Notr	Negatif
harika	*vardır*	*berbat*
mükemmel	*olmuştur*	*etmiyorum*
güzel	*tarafından*	*beğenmedim*
ederim	*bulunur*	*kötü*
memnunum	*bulunmaktadır*	*sıkıcı*
gayet	*değildir*	*iade*
1010	*sahiptir*	*rezalet*
iyi	*almıştır*	*en kötü*
teşekkürler	*aldı*	*yazık*
şık	*genellikle*	*kalitesiz*

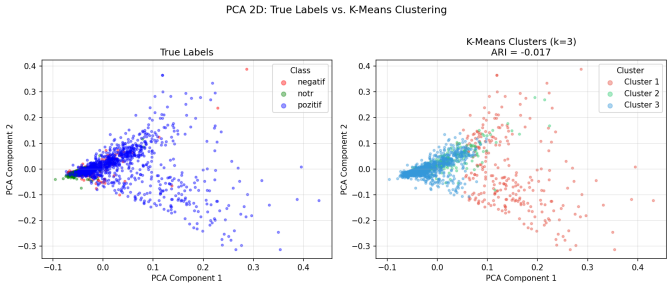

Fig. 5. PCA-2D: true labels (left) vs. K-Means ($k = 3$) clusters (right) on 2,000 test samples. ARI$= -0.017$ shows that sentiment classes do not form geometrically separable clusters in TF-IDF PCA space.

B. Classical Pattern Classifiers

I selected classifiers spanning different pattern recognition paradigms: probabilistic (Naive Bayes), maximum-margin (SVM), linear probabilistic (Logistic Regression), ensemble (Random Forest, XGBoost), instance-based (KNN), and single tree-based (Decision Tree) to cover the full spectrum from simple to complex decision boundaries. All use a TF-IDF pipeline with `norm='l2'`:

Naive Bayes (`MultinomialNB`, $\alpha = 0.5$): Models each class as a product of feature likelihoods. Simple and fast but assumes feature independence.

SVM (`LinearSVC`, $C = 0.5$, `max_iter=2000`): Finds the maximum-margin hyperplane. Well-suited to high-dimensional sparse TF-IDF features. Calibrated with `CalibratedClassifierCV` for probability output.

Logistic Regression (`saga` solver, $C = 0.5$): A probabilistic linear classifier. Competitive with SVM on text tasks and provides calibrated probabilities natively.

Random Forest (200 trees, `max_depth=None`): An ensemble of decision trees using bootstrap sampling and feature bagging. More expressive than linear models but slower on sparse TF-IDF.

KNN ($k = 5$, `metric=cosine`): Classifies by majority vote among k nearest neighbors in TF-IDF space. No explicit training phase; sensitive to document length normalization (handled by L2 norm).

Decision Tree (`max_depth=20`, `criterion=gini`): A single axis-aligned tree. Low computational cost but prone to overfitting in high-dimensional sparse feature spaces; included to demonstrate the capacity mismatch with TF-IDF.

XGBoost [12] (200 trees, `max_depth=5`, `learning_rate=0.05`): Gradient-boosted trees. Strong on structured data; on sparse TF-IDF it is slower than linear models and cannot leverage sparsity as effectively as SVM.

C. Transformer Classifier (BERTurk)

BERTurk (`dbmdz/bert-base-turkish-cased`) [5] is *fully fine-tuned* end-to-end for 3 epochs: all 12 transformer layers and the classification head are updated during training (no frozen layers, no adapters, no LoRA). Training used `learning_rate=2e-5`, `batch_size=32`, `max_length=128`, `AdamW` optimizer with `warmup_ratio=0.1` and `weight_decay=0.01`, on an NVIDIA RTX 4060 Laptop GPU (8 GB VRAM). BERTurk's WordPiece subword tokenizer handles Turkish morphological variation without manual stemming or lemmatization.

VI. Results

A. Evaluation Metrics

I report Accuracy, Macro F1, and Macro ROC-AUC (one-vs-rest). Macro F1 is the *primary* metric because the dataset is imbalanced. It penalizes models that ignore the minority *negatif* class equally to those that fail on the majority classes. ROC-AUC measures separability regardless of decision threshold.

B. Confusion Matrices and Negatif Class Analysis

Fig. 6 shows the confusion matrices for all eight models. The most important pattern recognition challenge is the low recall for the *negatif* class across all models due to severe class imbalance (11.6% of training data).

Table III quantifies this for the *negatif* class (support: 1,159 test samples). True Positives (TP), False Negatives (FN), False Positives (FP), and True Negatives (TN) are derived from the confusion matrix rows and columns corresponding to the negatif class.

The transformer recovers the most true negatives (TP=945, recall=81.5%), reflecting BERTurk's contextual understanding of negation. SVM achieves TP=757 (recall=65.3%), the best among classical models. Logistic Regression has the fewest false positives (FP=70) at the cost of very low recall (FN=563). Decision Tree has the highest false positive rate (FP=267) and the lowest recall alongside Random Forest, confirming that axis-aligned splits cannot effectively model the negatif minority class in high-dimensional TF-IDF space.

TABLE III
NEGATIF CLASS: TP, FN, FP, TN (TEST SET, $n = 10{,}000$)

Model	TP	FN	FP	TN
Transformer	945	214	151	8690
SVM	757	402	132	8709
Naive Bayes	675	484	168	8673
Logistic Reg.	596	563	70	8771
Random Forest	435	724	72	8769
KNN	477	682	132	8709
XGBoost	465	694	105	8736
Decision Tree	433	726	267	8574

C. Model Comparison

Table IV shows the main results for all eight models. BERTurk achieves the best Macro-F1 of 0.9298 and ROC-AUC of 0.9909. SVM is the best classical model (Macro-F1 = 0.8768, ROC-AUC = 0.9748), outperforming Naive Bayes by 2.8 F1 points and Logistic Regression by 3.8 points. XGBoost (Macro-F1 = 0.7554) and Decision Tree (Macro-F1 = 0.7063) are the weakest classical models: XGBoost's gradient boosting cannot efficiently exploit TF-IDF sparsity, and Decision Tree's axis-aligned splits collapse under high-dimensional feature pressure, yielding the largest gap between Accuracy (0.7935) and Macro-F1 (0.7063) of any model. The transformer-to-SVM gap is 5.3 Macro-F1 points, larger than on balanced datasets (see Appendix B), confirming that class imbalance amplifies the transformer's contextual advantage.

TABLE IV
TURKISH SENTIMENT CLASSIFICATION RESULTS ($n = 10{,}000$ TEST)

Model	Accuracy	Macro-F1	ROC-AUC
Transformer	**0.9574**	**0.9298**	**0.9909**
SVM	0.9223	0.8768	0.9748
Naive Bayes	0.9062	0.8488	0.9670
Logistic Reg.	0.9048	0.8387	0.9735
Random Forest	0.8863	0.7893	0.9625
KNN	0.8622	0.7757	0.9217
XGBoost	0.8369	0.7554	0.9257
Decision Tree	0.7935	0.7063	0.8282

Fig. 7 shows the model comparison bar chart across Accuracy, Macro-F1, Precision, and Recall. The consistent gap between Accuracy and Macro-F1 (4.5 points for SVM, 2.8 for Transformer) reflects the imbalance penalty. ROC curves are shown in Fig. 8.

D. Statistical Significance (10-Run Experiment)

To validate that SVM's advantage over other classical classifiers is statistically significant, I ran each model 10

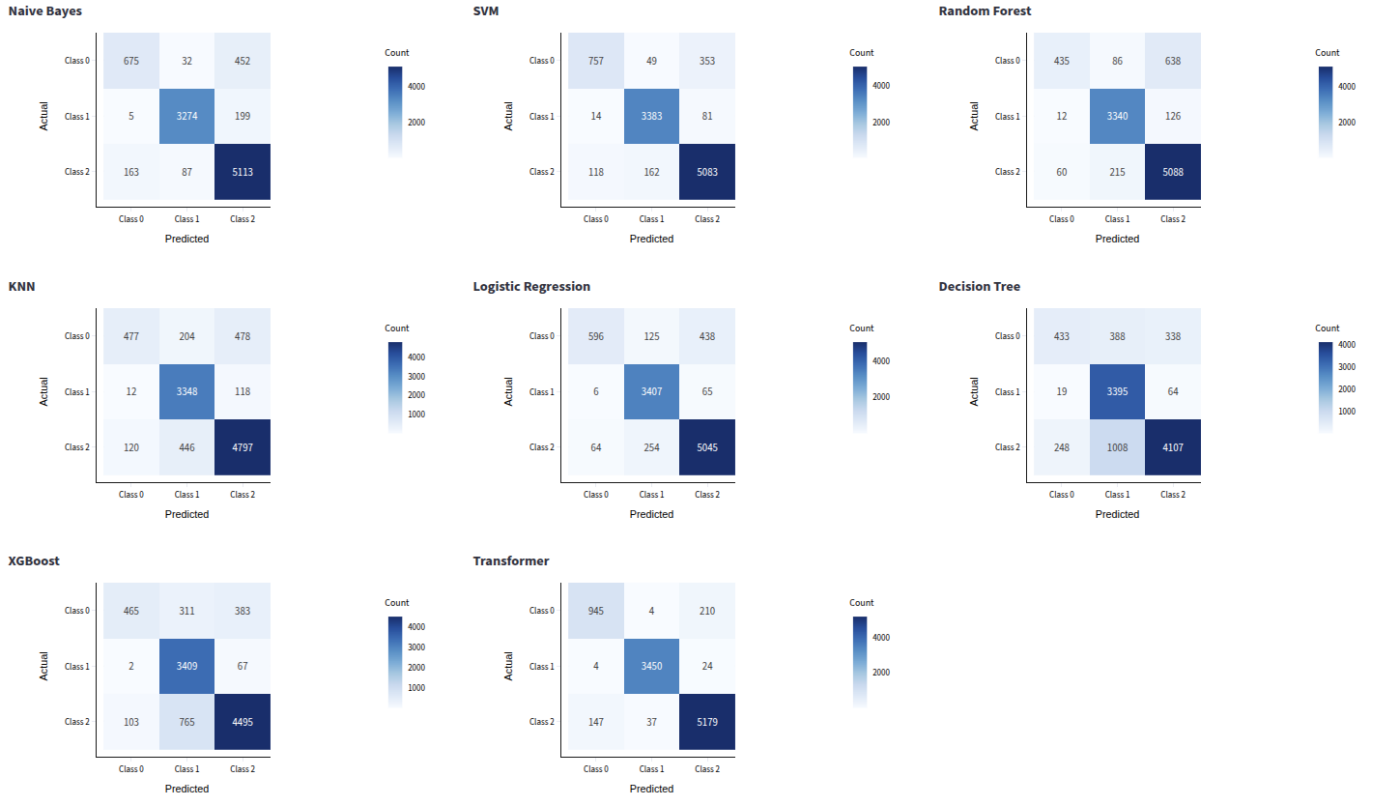

Fig. 6. Confusion matrices for all 8 models on the Turkish Sentiment test set ($n = 10{,}000$). Rows = true class, columns = predicted class. Classes: negatif (0), notr (1), pozitif (2). The negatif class (top row) shows the most misclassification across all models.

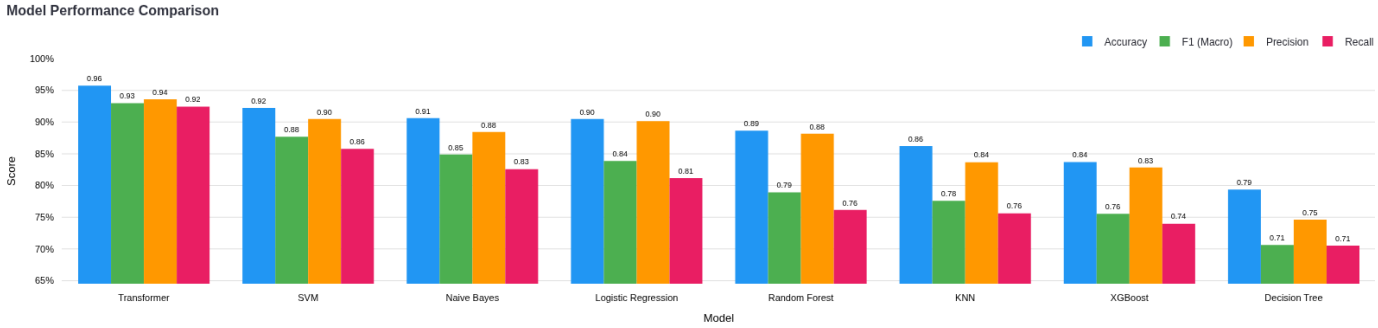

Fig. 7. Turkish Sentiment: Model performance comparison across Accuracy, Macro-F1, Precision, and Recall. The gap between Accuracy and Macro-F1 is largest for Random Forest (8.7 points) and smallest for Transformer (2.8 points).

times (seeds 0–9) on a stratified 10,000-sample subsample of the training set (80/20 internal split per run), using the same TF-IDF configuration. Table V reports mean Macro-F1, standard deviation, and two-sided independent t-test p-values comparing each model against SVM.

All p-values are below 0.0001, confirming that SVM significantly outperforms every other classical classifier across all seven models. Decision Tree shows the highest variance (std $= 0.0115$) and lowest mean F1 (0.675), consistent with its overfitting tendency in high-dimensional sparse spaces. XGBoost (0.754) outperforms Decision Tree and Random Forest but remains below the linear models, confirming that gradient boosting cannot efficiently exploit TF-IDF sparsity. SVM's advantage stems from three properties of LinearSVC: (1) it directly optimizes the margin in the high-dimensional TF-IDF space where linear separability is good; (2) it is not sensitive to feature independence assumptions (unlike Naive Bayes); and (3) L2 normalization stabilizes the margin even under the unequal class distribution. The transformer was excluded from the 10-run experiment due to GPU training

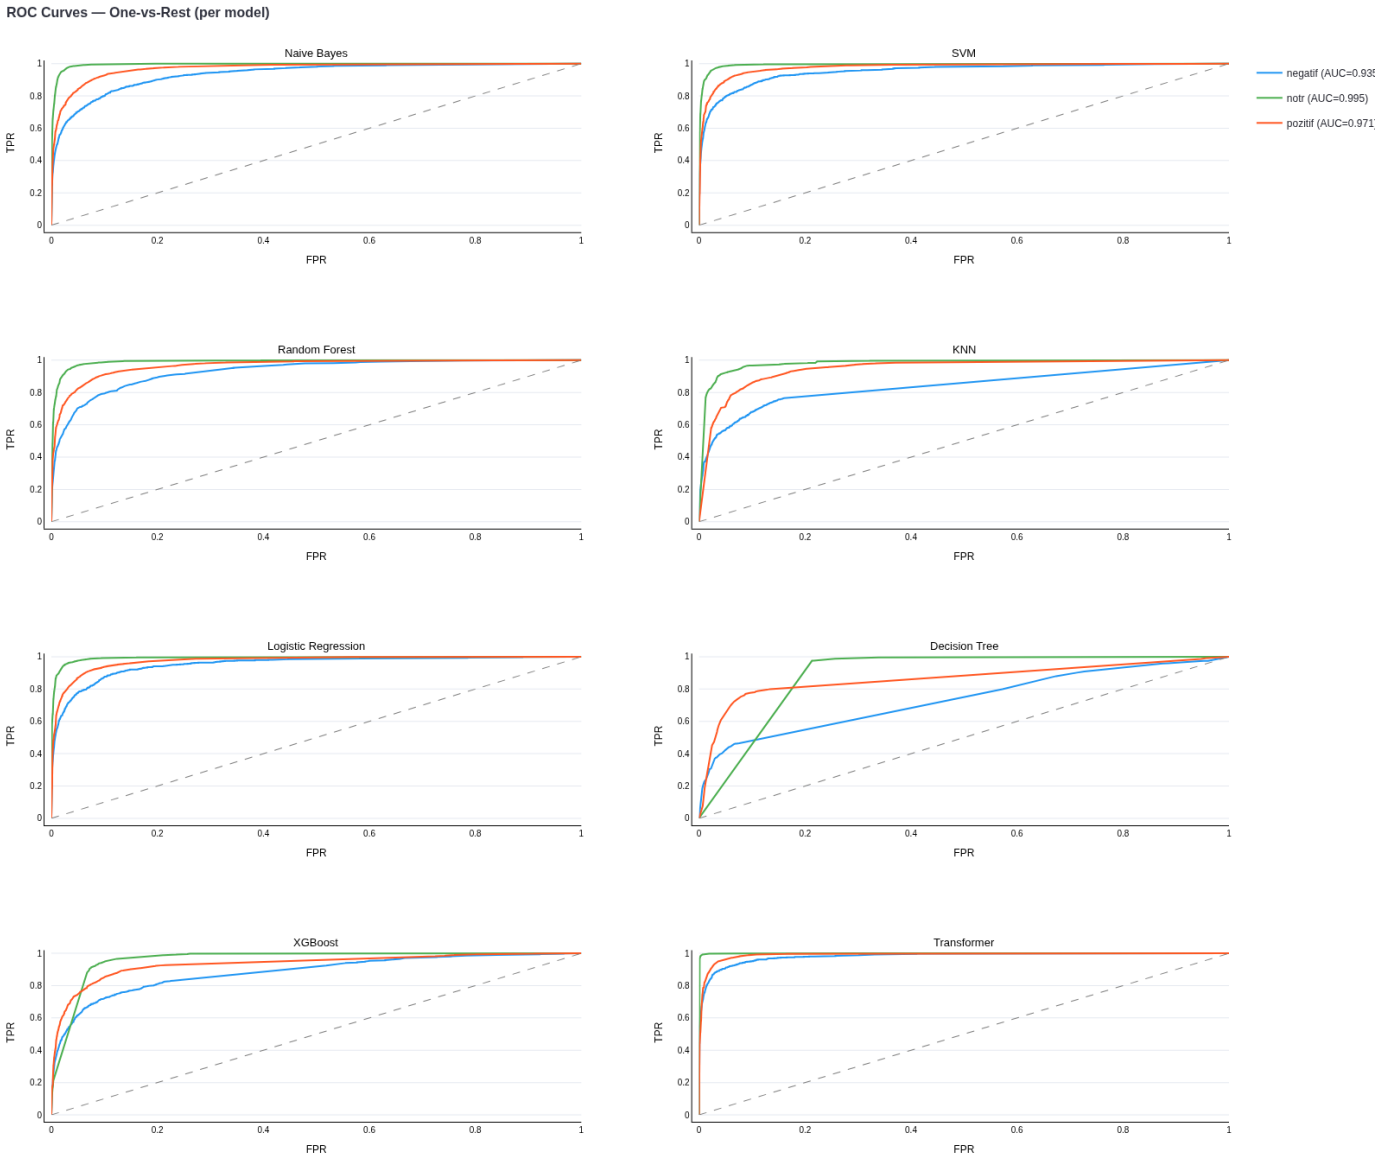

Fig. 8. Turkish Sentiment: ROC-AUC curves (one-vs-rest per class). Transformer macro AUC=0.991; SVM macro AUC=0.975. The negatif class (dashed lines) shows consistently lower AUC than notr and pozitif across all models.

TABLE V
10-RUN MACRO-F1: MEAN ± STD AND T-TEST VS. SVM

Model	Mean F1	Std	p-value vs. SVM
SVM	0.8214	0.0057	
Naive Bayes	0.8058	0.0070	<0.0001
Logistic Reg.	0.7733	0.0083	<0.0001
XGBoost	0.7539	0.0106	<0.0001
KNN	0.7413	0.0083	<0.0001
Random Forest	0.7360	0.0097	<0.0001
Decision Tree	0.6751	0.0115	<0.0001

cost (∼33 minutes per run).

VII. CONCLUSIONS

What was accomplished. I designed, trained, and evaluated a three-agent text classification system using classical pattern recognition methods and BERTurk on Turkish Sentiment Analysis. Across seven classical classifiers, SVM achieved the best classical Macro-F1 (0.877 on the test set; 0.821 mean over 10 runs) and BERTurk achieved the best overall result (Macro-F1 = 0.930). My 10-run statistical experiment confirmed SVM's superiority over all classical alternatives ($p < 0.0001$).

What was learned. The key pattern recognition insight is that high-dimensional sparse TF-IDF features provide a nearly linear feature space for sentiment classification, which favors

the maximum-margin approach of SVM. L2 normalization and linear decision boundaries in a 30K-dimensional space are more effective than non-linear ensemble methods (Random Forest, KNN) that cannot exploit global feature geometry. Class imbalance (negatif: 11.6%) is the primary challenge: even the transformer recovers only 81.5% of negatif instances. The K-Means ARI of -0.017 confirms that sentiment classes are not geometrically separable in PCA space, reinforcing that supervised boundary optimization (not cluster structure) drives performance.

What could not be done. Several things were outside my scope: (a) I did not apply strict 60/20/20 splits, using predefined benchmark splits instead; (b) I did not apply SMOTE or class-weighted loss to address negatif class under-representation; (c) I chose hyperparameters manually from literature recommendations rather than running a grid search; (d) I did not conduct quantitative evaluation of XAI faithfulness (explanation stability, fidelity metrics).

Future work. Applying SMOTE or class-weighted LinearSVC to improve negatif recall is the most impactful next step. Grid-search over SVM's C parameter and TF-IDF vocabulary size could close part of the transformer gap. Quantitative XAI faithfulness evaluation (e.g., sufficiency and comprehensiveness metrics) would complete the pattern recognition analysis beyond classification accuracy.

References

[1] R. O. Duda, P. E. Hart, and D. G. Stork, *Pattern Classification*, 2nd ed. New York, NY, USA: Wiley, 2001.

[2] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of deep bidirectional transformers for language understanding," in *Proc. NAACL-HLT*, Minneapolis, MN, USA, Jun. 2019, pp. 4171–4186.

[3] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Proc. NeurIPS*, Long Beach, CA, USA, Dec. 2017, pp. 4765–4774.

[4] M. T. Ribeiro, S. Singh, and C. Guestrin, "'Why should I trust you?': Explaining the predictions of any classifier," in *Proc. KDD*, San Francisco, CA, USA, Aug. 2016, pp. 1135–1144.

[5] S. Schweter, "BERTurk: BERT models for Turkish," *Zenodo*, Apr. 2020. doi: 10.5281/zenodo.3770924.

[6] S. Minaee et al., "Deep learning–based text classification: A comprehensive review," *ACM Comput. Surv.*, vol. 54, no. 3, pp. 1–40, Apr. 2021.

[7] M. Wooldridge and N. R. Jennings, "Intelligent agents: Theory and practice," *Knowl. Eng. Rev.*, vol. 10, no. 2, pp. 115–152, Jun. 1995.

[8] winvoker, "Turkish Sentiment Analysis Dataset," Hugging Face Datasets, 2022. [Online]. Available: https://huggingface.co/datasets/winvoker/turkish-sentiment-analysis-dataset

[9] stanfordnlp, "Large Movie Review Dataset (IMDB)," Hugging Face Datasets, 2022. [Online]. Available: https://huggingface.co/datasets/stanfordnlp/imdb

[10] sh0416, "AG's News Topic Classification Dataset," Hugging Face Datasets, 2023. [Online]. Available: https://huggingface.co/datasets/sh0416/ag_news

[11] S. Aras, "TTC4900: A Benchmark Data for Turkish Text Categorization," Hugging Face Datasets, 2022. [Online]. Available: https://huggingface.co/datasets/savasy/ttc4900

[12] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proc. KDD*, San Francisco, CA, USA, Aug. 2016, pp. 785–794.

[13] L. van der Maaten and G. Hinton, "Visualizing data using t-SNE," *J. Mach. Learn. Res.*, vol. 9, pp. 2579–2605, Nov. 2008.

[14] G. Salton and C. Buckley, "Term-weighting approaches in automatic text retrieval," *Inf. Process. Manag.*, vol. 24, no. 5, pp. 513–523, 1988.

Appendix

The full source code, dataset configuration files, experiment scripts, and all figures for this project are available at:

https://github.com/oruccakir/multi-agent-xai-text-classifier

The repository is organised as follows: `src/` contains training and evaluation pipelines; `agents/` contains the three-agent routing logic; `app/` contains the Streamlit interface; `reports/` contains this document and all generated figures; `scripts/` contains figure generation scripts for this report.

Table VI summarises results for the three additional datasets used in this project. Full confusion matrices, ROC curves, and per-model figures are available in the GitHub repository.

TABLE VI
ADDITIONAL DATASETS: BEST CLASSICAL MODEL AND TRANSFORMER MACRO-F1

Dataset	Task	Best Classical	F1-M	Transformer F1-M
IMDB	Bin. Sentiment	SVM	0.895	0.907
AG News	4-class Topics	SVM	0.915	0.928
Turkish News	7-class Topics	SVM	0.919	0.938

Cross-dataset patterns. SVM is the best classical model on all three datasets. The transformer-to-SVM gap is 1.2–1.9 F1 points on balanced datasets, much smaller than the 5.3-point gap on Turkish Sentiment. This confirms that class imbalance amplifies the transformer's advantage. Decision Tree collapses on multi-class tasks (AG News F1=0.611, Turkish News F1=0.491), while KNN performs competitively on short text (Turkish News F1=0.884).